

urn-prediction-logistic-regression

May 25, 2024

```
[33]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import KFold, StratifiedKFold, train_test_split
from sklearn.metrics import roc_auc_score, accuracy_score, confusion_matrix, \
    roc_curve, precision_score, recall_score, precision_recall_curve
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=UserWarning)
```

```
[2]: df = pd.read_csv('/content/churn_prediction.csv')
```

```
[3]: df.head()
```

```
[3]:
```

	customer_id	vintage	age	gender	dependents	occupation	city	\
0	1	3135	66	Male	0.0	self_employed	187.0	
1	2	310	35	Male	0.0	self_employed	NaN	
2	4	2356	31	Male	0.0	salaried	146.0	
3	5	478	90	NaN	NaN	self_employed	1020.0	
4	6	2531	42	Male	2.0	self_employed	1494.0	

	customer_nw_category	branch_code	days_since_last_transaction	...	\
0	2	755	224.0	...	
1	2	3214	60.0	...	
2	2	41	NaN	...	
3	2	582	147.0	...	
4	3	388	58.0	...	

	previous_month_end_balance	average_monthly_balance_prevQ	\
0	1458.71	1458.71	
1	8704.66	7799.26	
2	5815.29	4910.17	
3	2291.91	2084.54	

4	1401.72	1643.31
---	---------	---------

	average_monthly_balance_prevQ2	current_month_credit \
0	1449.07	0.20
1	12419.41	0.56
2	2815.94	0.61
3	1006.54	0.47
4	1871.12	0.33

	previous_month_credit	current_month_debit	previous_month_debit \
0	0.20	0.20	0.20
1	0.56	5486.27	100.56
2	0.61	6046.73	259.23
3	0.47	0.47	2143.33
4	714.61	588.62	1538.06

	current_month_balance	previous_month_balance	churn
0	1458.71	1458.71	0
1	6496.78	8787.61	0
2	5006.28	5070.14	0
3	2291.91	1669.79	1
4	1157.15	1677.16	1

[5 rows x 21 columns]

```
[4]: df.columns
```

```
[4]: Index(['customer_id', 'vintage', 'age', 'gender', 'dependents', 'occupation',
'city', 'customer_nw_category', 'branch_code',
'days_since_last_transaction', 'current_balance',
'previous_month_end_balance', 'average_monthly_balance_prevQ',
'average_monthly_balance_prevQ2', 'current_month_credit',
'previous_month_credit', 'current_month_debit', 'previous_month_debit',
'current_month_balance', 'previous_month_balance', 'churn'],
dtype='object')
```

```
[5]: df.describe()
```

```
[5]:
```

	customer_id	vintage	age	dependents	city \
count	28382.000000	28382.000000	28382.000000	25919.000000	27579.000000
mean	15143.508667	2364.336446	48.208336	0.347236	796.109576
std	8746.454456	1610.124506	17.807163	0.997661	432.872102
min	1.000000	180.000000	1.000000	0.000000	0.000000
25%	7557.250000	1121.000000	36.000000	0.000000	409.000000
50%	15150.500000	2018.000000	46.000000	0.000000	834.000000
75%	22706.750000	3176.000000	60.000000	0.000000	1096.000000
max	30301.000000	12899.000000	90.000000	52.000000	1649.000000

	customer_nw_category	branch_code	days_since_last_transaction \
count	28382.000000	28382.000000	25159.000000
mean	2.225530	925.975019	69.997814
std	0.660443	937.799129	86.341098
min	1.000000	1.000000	0.000000
25%	2.000000	176.000000	11.000000
50%	2.000000	572.000000	30.000000
75%	3.000000	1440.000000	95.000000
max	3.000000	4782.000000	365.000000

	current_balance	previous_month_end_balance \
count	2.838200e+04	2.838200e+04
mean	7.380552e+03	7.495771e+03
std	4.259871e+04	4.252935e+04
min	-5.503960e+03	-3.149570e+03
25%	1.784470e+03	1.906000e+03
50%	3.281255e+03	3.379915e+03
75%	6.635820e+03	6.656535e+03
max	5.905904e+06	5.740439e+06

	average_monthly_balance_prevQ	average_monthly_balance_prevQ2 \
count	2.838200e+04	2.838200e+04
mean	7.496780e+03	7.124209e+03
std	4.172622e+04	4.457581e+04
min	1.428690e+03	-1.650610e+04
25%	2.180945e+03	1.832507e+03
50%	3.542865e+03	3.359600e+03
75%	6.666887e+03	6.517960e+03
max	5.700290e+06	5.010170e+06

	current_month_credit	previous_month_credit	current_month_debit \
count	2.838200e+04	2.838200e+04	2.838200e+04
mean	3.433252e+03	3.261694e+03	3.658745e+03
std	7.707145e+04	2.968889e+04	5.198542e+04
min	1.000000e-02	1.000000e-02	1.000000e-02
25%	3.100000e-01	3.300000e-01	4.100000e-01
50%	6.100000e-01	6.300000e-01	9.193000e+01
75%	7.072725e+02	7.492350e+02	1.360435e+03
max	1.226985e+07	2.361808e+06	7.637857e+06

	previous_month_debit	current_month_balance	previous_month_balance \
count	2.838200e+04	2.838200e+04	2.838200e+04
mean	3.339761e+03	7.451133e+03	7.495177e+03
std	2.430111e+04	4.203394e+04	4.243198e+04
min	1.000000e-02	-3.374180e+03	-5.171920e+03
25%	4.100000e-01	1.996765e+03	2.074407e+03

50%	1.099600e+02	3.447995e+03	3.465235e+03
75%	1.357553e+03	6.667958e+03	6.654693e+03
max	1.414168e+06	5.778185e+06	5.720144e+06

```

churn
count    28382.000000
mean      0.185329
std       0.388571
min       0.000000
25%      0.000000
50%      0.000000
75%      0.000000
max       1.000000

```

```
[6]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28382 entries, 0 to 28381
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   customer_id                          28382 non-null  int64
1   vintage                              28382 non-null  int64
2   age                                  28382 non-null  int64
3   gender                               27857 non-null  object
4   dependents                           25919 non-null  float64
5   occupation                           28302 non-null  object
6   city                                  27579 non-null  float64
7   customer_nw_category                 28382 non-null  int64
8   branch_code                          28382 non-null  int64
9   days_since_last_transaction           25159 non-null  float64
10  current_balance                       28382 non-null  float64
11  previous_month_end_balance             28382 non-null  float64
12  average_monthly_balance_prevQ         28382 non-null  float64
13  average_monthly_balance_prevQ2       28382 non-null  float64
14  current_month_credit                   28382 non-null  float64
15  previous_month_credit                  28382 non-null  float64
16  current_month_debit                    28382 non-null  float64
17  previous_month_debit                   28382 non-null  float64
18  current_month_balance                  28382 non-null  float64
19  previous_month_balance                  28382 non-null  float64
20  churn                                  28382 non-null  int64
dtypes: float64(13), int64(6), object(2)
memory usage: 4.5+ MB

```

0.0.1 Handling Missing Values

```
[7]: df.isnull().sum()
```

```
[7]: customer_id          0
      vintage            0
      age               0
      gender            525
      dependents       2463
      occupation        80
      city             803
      customer_nw_category  0
      branch_code       0
      days_since_last_transaction 3223
      current_balance    0
      previous_month_end_balance  0
      average_monthly_balance_prevQ  0
      average_monthly_balance_prevQ2  0
      current_month_credit  0
      previous_month_credit  0
      current_month_debit  0
      previous_month_debit  0
      current_month_balance  0
      previous_month_balance  0
      churn              0
      dtype: int64
```

```
[8]: df['gender'].replace({'Male':1, 'Female':0}, inplace=True)
```

```
[9]: df['gender'].value_counts()
```

```
[9]: gender
      1.0    16548
      0.0    11309
      Name: count, dtype: int64
```

```
[10]: df['gender']
```

```
[10]: 0      1.0
      1      1.0
      2      1.0
      3      NaN
      4      1.0
      ...
      28377  0.0
      28378  0.0
      28379  1.0
```

```
28380    1.0
28381    1.0
Name: gender, Length: 28382, dtype: float64
```

```
[11]: mode_gender = df['gender'].mode()[0]
```

```
[12]: df['gender'].fillna(mode_gender,inplace=True)
```

```
[13]: df['gender']
```

```
[13]: 0      1.0
      1      1.0
      2      1.0
      3      1.0
      4      1.0
      ...
      28377  0.0
      28378  0.0
      28379  1.0
      28380  1.0
      28381  1.0
Name: gender, Length: 28382, dtype: float64
```

```
[14]: df['gender'].value_counts()
```

```
[14]: gender
      1.0    17073
      0.0    11309
Name: count, dtype: int64
```

```
[15]: df['dependents'].value_counts()
```

```
[15]: dependents
      0.0    21435
      2.0    2150
      1.0    1395
      3.0     701
      4.0    179
      5.0     41
      6.0      8
      7.0      3
      9.0      1
      52.0     1
      36.0     1
      50.0     1
      8.0      1
      25.0     1
```

```
32.0      1
Name: count, dtype: int64
```

```
[16]: df['dependents'] = df['dependents'].fillna(0)
```

```
[17]: df['dependents'].value_counts()
```

```
[17]: dependents
0.0      23898
2.0       2150
1.0       1395
3.0        701
4.0        179
5.0         41
6.0          8
7.0          3
9.0          1
52.0         1
36.0         1
50.0         1
8.0          1
25.0         1
32.0         1
Name: count, dtype: int64
```

```
[18]: df['occupation'].fillna('self_employed',inplace=True)
```

```
[19]: df['city'].mode()
```

```
[19]: 0      1020.0
Name: city, dtype: float64
```

```
[20]: df['city'].fillna(df['city'].mode()[0],inplace=True)
```

```
[21]: df['city'].isnull().sum()
```

```
[21]: 0
```

```
[22]: df['days_since_last_transaction'].fillna(df['days_since_last_transaction'].
↪mode()[0],inplace=True)
```

```
[23]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28382 entries, 0 to 28381
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype

```

```

---  -----
0   customer_id          28382 non-null  int64
1   vintage              28382 non-null  int64
2   age                  28382 non-null  int64
3   gender                28382 non-null  float64
4   dependents           28382 non-null  float64
5   occupation           28382 non-null  object
6   city                 28382 non-null  float64
7   customer_nw_category 28382 non-null  int64
8   branch_code          28382 non-null  int64
9   days_since_last_transaction 28382 non-null  float64
10  current_balance      28382 non-null  float64
11  previous_month_end_balance 28382 non-null  float64
12  average_monthly_balance_prevQ 28382 non-null  float64
13  average_monthly_balance_prevQ2 28382 non-null  float64
14  current_month_credit 28382 non-null  float64
15  previous_month_credit 28382 non-null  float64
16  current_month_debit  28382 non-null  float64
17  previous_month_debit 28382 non-null  float64
18  current_month_balance 28382 non-null  float64
19  previous_month_balance 28382 non-null  float64
20  churn                28382 non-null  int64

```

dtypes: float64(14), int64(6), object(1)

memory usage: 4.5+ MB

```
[24]: df.isnull().sum()
```

```

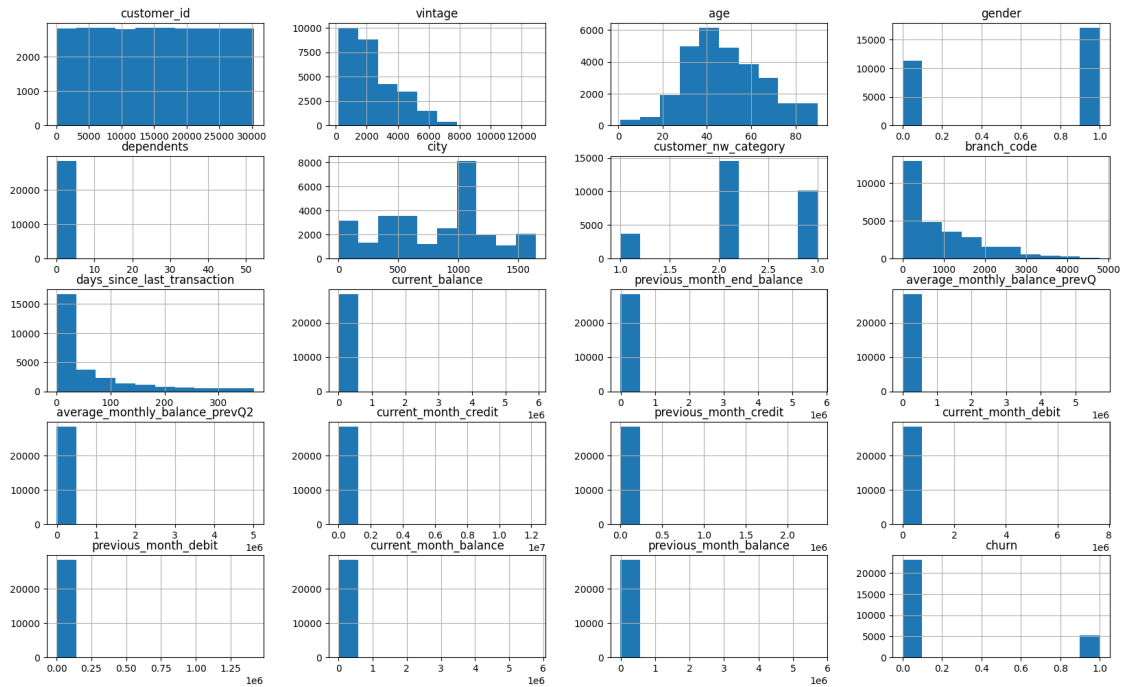
[24]: customer_id          0
      vintage              0
      age                  0
      gender                0
      dependents           0
      occupation           0
      city                 0
      customer_nw_category 0
      branch_code          0
      days_since_last_transaction 0
      current_balance      0
      previous_month_end_balance 0
      average_monthly_balance_prevQ 0
      average_monthly_balance_prevQ2 0
      current_month_credit 0
      previous_month_credit 0
      current_month_debit  0
      previous_month_debit 0
      current_month_balance 0
      previous_month_balance 0

```



churn
dtype: int64

0


```
[25]: df.hist(figsize=(20, 12))  
plt.show()
```



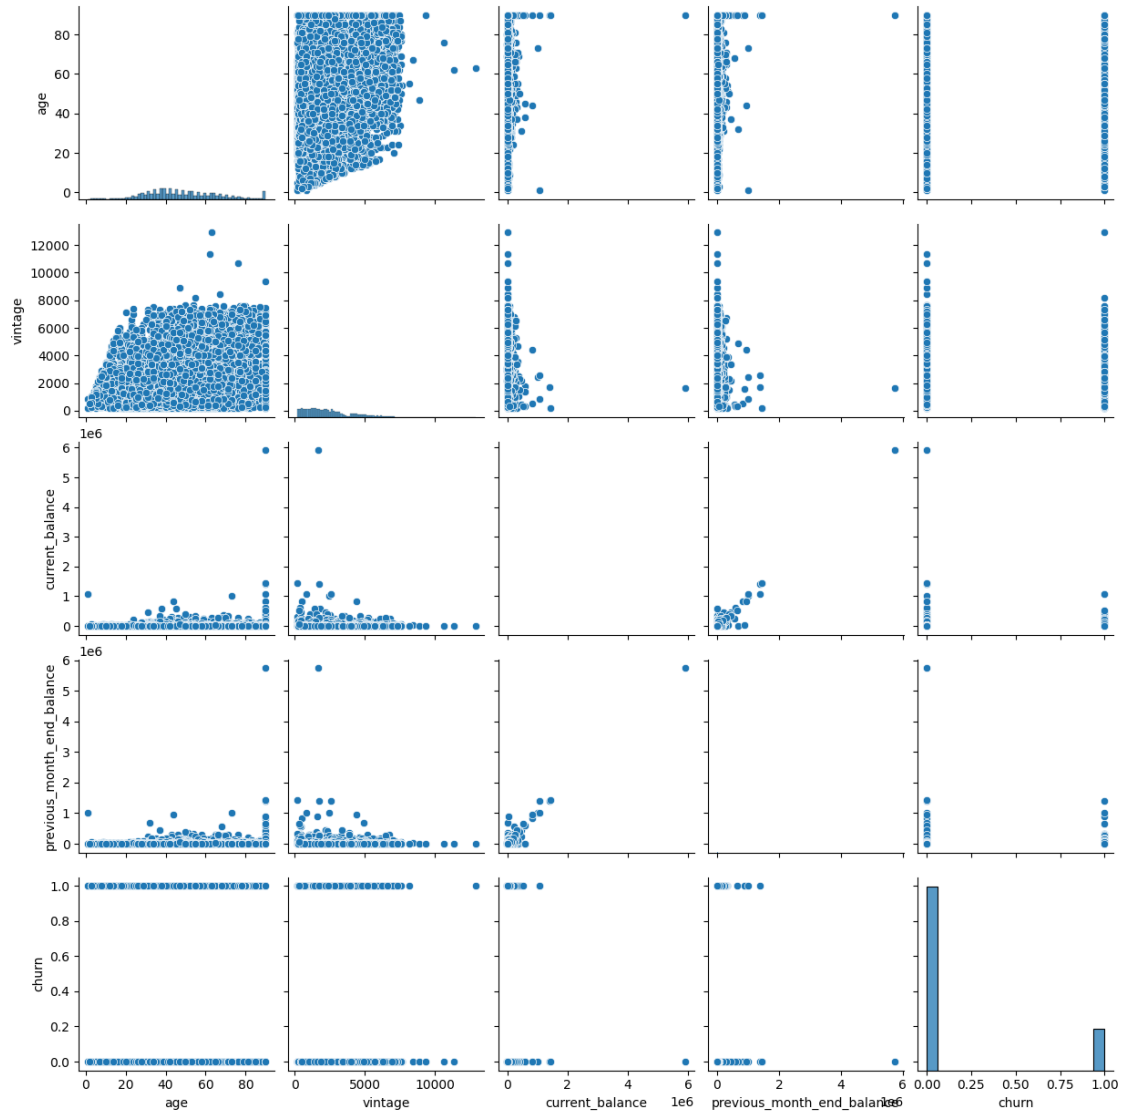
```
[26]: sns.countplot(data=df, x='occupation')  
plt.xticks(rotation=90)  
plt.show()
```




```
[27]: sns.countplot(data=df, x='gender')  
plt.show()
```




```
[28]: sns.pairplot(df[['age', 'vintage', 'current_balance',  
    ↪ 'previous_month_end_balance', 'churn']])  
plt.show()
```




```
[29]: sns.histplot(df['vintage'], kde=True)
plt.show()
```



```
[30]: sns.countplot(data=df, x='churn')  
plt.show()
```



```
[31]: sns.histplot(df['days_since_last_transaction'], kde=True)  
plt.show()
```



```
[32]: df = pd.concat([df,pd.get_dummies(df['occupation'],prefix =_
↳str('occupation'),prefix_sep='_')],axis = 1)
```

```
[34]: num_cols = ['customer_nw_category', 'current_balance',
                 'previous_month_end_balance', 'average_monthly_balance_prevQ2',_
↳'average_monthly_balance_prevQ',
                 'current_month_credit','previous_month_credit',_
↳'current_month_debit',
                 'previous_month_debit','current_month_balance',_
↳'previous_month_balance']
for i in num_cols:
    df[i] = np.log(df[i] + 17000)

std = StandardScaler()
scaled = std.fit_transform(df[num_cols])
scaled = pd.DataFrame(scaled,columns=num_cols)
```

```
[35]: df_df Og = df.copy()
df = df.drop(columns = num_cols,axis = 1)
df = df.merge(scaled,left_index=True,right_index=True,how = "left")
```


```
[36]: y_all = df.churn
df = df.drop(['churn','customer_id','occupation'],axis = 1)

[37]: baseline_cols = ['current_month_debit',□
↳'previous_month_debit','current_balance','previous_month_end_balance','vintage'
, 'occupation_retired',□
↳'occupation_salaried','occupation_self-employed', 'occupation_student']

[39]: df_baseline = df[baseline_cols]
xtrain, xtest, ytrain, ytest = train_test_split(df_baseline,y_all,test_size=1/
↳3, random_state=11, stratify = y_all)

[40]: model = LogisticRegression()
model.fit(xtrain,ytrain)
pred = model.predict_proba(xtest)[:,-1]

[41]: from sklearn.metrics import roc_curve
fpr, tpr, _ = roc_curve(ytest,pred)
auc = roc_auc_score(ytest, pred)
plt.figure(figsize=(12,8))
plt.plot(fpr,tpr,label="Validation AUC-ROC="+str(auc))
x = np.linspace(0, 1, 1000)
plt.plot(x, x, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc=4)
plt.show()
```





```
[42]: pred_val = model.predict(xtest)
```

```
[43]: label_preds = pred_val

cm = confusion_matrix(ytest,label_preds)

def plot_confusion_matrix(cm, normalized=True, cmap='bone'):
    plt.figure(figsize=[7, 6])
    norm_cm = cm
    if normalized:
        norm_cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        sns.heatmap(norm_cm, annot=cm, fmt='g', xticklabels=['Predicted: No', 'Predicted: Yes'], yticklabels=['Actual: No', 'Actual: Yes'], cmap=cmap)
    plot_confusion_matrix(cm, ['No', 'Yes'])
```



```
[44]: recall_score(ytest,pred_val)
```

```
[44]: 0.11580148317170565
```

```
[45]: def cv_score(ml_model, rstate = 12, thres = 0.5, cols = df.columns):
    i = 1
    cv_scores = []
    df1 = df.copy()
    df1 = df[cols]

    # 5 Fold cross validation stratified on the basis of target
    kf = StratifiedKFold(n_splits=5,random_state=rstate,shuffle=True)
    for df_index,test_index in kf.split(df1,y_all):
        print('\n{} of kfold {}'.format(i,kf.n_splits))
        xtr,xvl = df1.loc[df_index],df1.loc[test_index]
        ytr,yvl = y_all.loc[df_index],y_all.loc[test_index]
```

```

# Define model for fitting on the training set for each fold
model = ml_model
model.fit(xtr, ytr)
pred_probs = model.predict_proba(xvl)
pp = []

# Use threshold to define the classes based on probability values
for j in pred_probs[:,1]:
    if j>thres:
        pp.append(1)
    else:
        pp.append(0)

# Calculate scores for each fold and print
pred_val = pp
roc_score = roc_auc_score(yvl,pred_probs[:,1])
recall = recall_score(yvl,pred_val)
precision = precision_score(yvl,pred_val)
suffix = ""
msg = ""
msg += "ROC AUC Score: {}, Recall Score: {:.4f}, Precision Score: {:.
↪4f} ".format(roc_score, recall,precision)
print("{}".format(msg))

# Save scores
cv_scores.append(roc_score)
i+=1
return cv_scores

```

```
[46]: baseline_scores = cv_score(LogisticRegression(), cols = baseline_cols)
```

1 of kfold 5

ROC AUC Score: 0.7644836090843695, Recall Score: 0.0751, Precision Score: 0.5766

2 of kfold 5

ROC AUC Score: 0.7805820573425136, Recall Score: 0.0722, Precision Score: 0.6552

3 of kfold 5

ROC AUC Score: 0.755325784138303, Recall Score: 0.1350, Precision Score: 0.6455

4 of kfold 5

ROC AUC Score: 0.7582064809820148, Recall Score: 0.1169, Precision Score: 0.6508

5 of kfold 5

ROC AUC Score: 0.7622577114607865, Recall Score: 0.1112, Precision Score: 0.5821

```
[47]: all_feat_scores = cv_score(LogisticRegression())
```

```
1 of kfold 5  
ROC AUC Score: 0.7445534888500668, Recall Score: 0.1274, Precision Score: 0.7016
```

```
2 of kfold 5  
ROC AUC Score: 0.7296795807214058, Recall Score: 0.1188, Precision Score: 0.6281
```

```
3 of kfold 5  
ROC AUC Score: 0.7184804935729603, Recall Score: 0.1131, Precision Score: 0.5667
```

```
4 of kfold 5  
ROC AUC Score: 0.758577951701817, Recall Score: 0.2015, Precision Score: 0.6795
```

```
5 of kfold 5  
ROC AUC Score: 0.7115271866407042, Recall Score: 0.0970, Precision Score: 0.4928
```

```
[48]: from sklearn.feature_selection import RFE  
import matplotlib.pyplot as plt  
  
model = LogisticRegression()  
rfe = RFE(estimator=model, n_features_to_select=1, step=1)  
rfe.fit(df, y_all)
```

```
[48]: RFE(estimator=LogisticRegression(), n_features_to_select=1)
```

```
[49]: ranking_df = pd.DataFrame()  
ranking_df['Feature_name'] = df.columns  
ranking_df['Rank'] = rfe.ranking_
```

```
[50]: ranked = ranking_df.sort_values(by=['Rank'])
```

```
[51]: rfe_top_10_scores = cv_score(LogisticRegression(), cols =  
    ↪ranked['Feature_name'][:10].values)
```

```
1 of kfold 5  
ROC AUC Score: 0.7999459459459459, Recall Score: 0.2310, Precision Score: 0.7254
```

```
2 of kfold 5  
ROC AUC Score: 0.8076882129277567, Recall Score: 0.2224, Precision Score: 0.7548
```

```
3 of kfold 5  
ROC AUC Score: 0.7994672776849501, Recall Score: 0.2253, Precision Score: 0.7096
```

```
4 of kfold 5  
ROC AUC Score: 0.7944628044127514, Recall Score: 0.2063, Precision Score: 0.7282
```

```
5 of kfold 5
ROC AUC Score: 0.7950266916205085, Recall Score: 0.1920, Precision Score: 0.6733
```

```
[52]: cv_score(LogisticRegression(), cols = ranked['Feature_name'][:10].values,
          ↪thres=0.14)
```

```
1 of kfold 5
ROC AUC Score: 0.7999459459459459, Recall Score: 0.8327, Precision Score: 0.2888
```

```
2 of kfold 5
ROC AUC Score: 0.8076882129277567, Recall Score: 0.8365, Precision Score: 0.2952
```

```
3 of kfold 5
ROC AUC Score: 0.7994672776849501, Recall Score: 0.8241, Precision Score: 0.2951
```

```
4 of kfold 5
ROC AUC Score: 0.7944628044127514, Recall Score: 0.8194, Precision Score: 0.2882
```


```
5 of kfold 5
ROC AUC Score: 0.7950266916205085, Recall Score: 0.8099, Precision Score: 0.2993
```

```
[52]: [0.7999459459459459,
      0.8076882129277567,
      0.7994672776849501,
      0.7944628044127514,
      0.7950266916205085]
```

```
[53]: results_df = pd.DataFrame({'baseline':baseline_scores, 'all_feats':
          ↪all_feat_scores, 'rfe_top_10': rfe_top_10_scores})
```

```
[54]: results_df.plot(y=["baseline", "all_feats", "rfe_top_10"], kind="bar")
```

```
[54]: <Axes: >
```



0.0.2 Here, we can see that the model based on RFE is giving the best result for each fold.

[]: